The lab Flat index compares a query vector with every stored vector and returns the smallest L2 distances.

The lab HNSW index uses a graph. Each node keeps 32 neighbors. Construction uses efConstruction 40.

The lab IVF index splits vectors into 100 lists. A query searches 10 of those lists.

BM25 reranking scores a retrieved chunk by its row in the full corpus. A higher BM25 score is ordered first.

Recall counts a retrieved chunk as relevant when its document and page were labeled in the golden file.

The public demo does not call OpenAI. Local generation sends the retrieved chunk text to the configured OpenAI model.